

KESHAV GARG

Delhi, India | +91 9873501040 | gargkeshav504@gmail.com | LinkedIn | Portfolio

EXPERIENCE

AI Engineer

Apr 2025 – Present

Kroolo

Delhi, India

- Architected and built a complete Enterprise Search system from the ground up, managing the full backend lifecycle—from data ingestion and preprocessing to a search pipeline using traditional machine learning methods—achieving performance comparable to an industry leader.
- Architected the Enterprise Search platform integrating 10+ apps with search and multi-agent orchestration.
- Revamped Chat with Anything module using advanced text extraction algorithms, reducing processing latency by 50%.
- Rebuilt AI agent ecosystem with 250+ integrated tools, transforming prompt-engineered agents into sophisticated agentic systems.
- Performed SFT fine-tuning of FunctionGemma for local agentic use, enabling on-device tool-calling with reduced inference cost and improved instruction-following.
- Led technical design reviews and managed full project lifecycle from requirements to production deployment.
- Delivered comprehensive system documentation for enterprise-grade AI agent platform.

AI Engineer (Contract)

March 2025 – Present

GenAI Protos

Delhi, India

- Developed and deployed privacy-focused AI applications, including an Enterprise Search engine and Legal Assistant, using local LLM inference on NVIDIA DGX Spark.
- Engineered a multimodal Deep Research Multi-Agent System with 97% accuracy for fact-checking and misinformation detection.
- Developed NLP2SQL system using RAG architecture, achieving 98% query accuracy with automated chart generation, deployed via WhatsApp.
- Optimized open-source LLMs on NVIDIA DGX Spark for high-throughput, low-latency inferencing on on-premises enterprise deployments under strict security constraints.
- Built YOLO-based Object Detection Model with custom fine-tuning, eliminating 100% manual CCTV monitoring.

Freelance AI Engineer

Nov 2025 – Present

Independent — Enterprise SaaS Client

Remote

- Designed and shipped two AI products for an enterprise SaaS client: a Pricing platform for AI-assisted proposal authoring and a Pricing-IGCE module for automated cost-estimate generation, reducing manual proposal effort by 95%.
- Built multi-agent workflows that parse inbound RFPs, draft technical narratives, compute labor and indirect rates, and export priced proposals to Excel.
- Owned product end-to-end across frontend, backend, agent orchestration, and deployment.

AI/ML Intern

May 2024 – July 2024

Agnisys Technology

Noida, India

- Implemented a Retrieval-Augmented Generation (RAG) pipeline, significantly improving data retrieval accuracy and contextual response generation.
- Created an automated dashboard using user behavior data, eliminating manual data entry and reducing processing time.
- Collaborated with cross-functional teams to deliver scalable, high-impact solutions, improving project outcomes and client satisfaction.

PROJECTS

Ra.One – Human-like AI WhatsApp Assistant | Demo

- Built an AI-powered WhatsApp assistant on LangGraph with dual-memory architecture (PostgreSQL short-term + Pinecone long-term) and multimodal inputs including ElevenLabs voice.

Enterprise Support Automation | Demo

- Autonomous LangGraph agent that scrapes 100+ help articles via Selenium, embeds them for retrieval, and resolves user queries with verification — reducing manual workload by 80%.

LegalInsight | Demo

- Legal document analysis platform using RAPTOR for hierarchical understanding and “Chat with PDF” retrieval — 70% faster review — with privacy-first deployment via local Ollama LLMs and PyMuPDF4LLM extraction.

EDUCATION

Bachelor of Technology (Computer Science)

2021 – 2025

Amity University Noida

Noida, India

TECHNICAL SKILLS

Programming: Python, SQL, TypeScript

Machine Learning & AI: Large Language Models (LLMs), Retrieval-Augmented Generation (RAG), Multi-Agent Systems, Fine-tuning, Deep Learning, Natural Language Processing (NLP)

ML Frameworks: PyTorch, Scikit-learn, Pandas, LangChain, LangGraph, Agno, Hugging Face

Backend & Infra: FastAPI, Next.js, AWS, Docker, Railway, Stripe, PostgreSQL